

ComicCamp: A Framework for Comic-Style Campaign Narration

Ashutosh Ranjan, Vivek Srivastava, Shirish Karande
TCS Research, India

{ashutosh.ranjan2, srivastava.vivek2, shirish.karande}@tcs.com

1 Introduction

Email broadcast campaigns are powerful tools to engage with the target audience [7, 17] to broadcast messages such as corporate announcement, policy update, webinar invitation, etc. The campaign creators often look out for creative ways to compose and communicate such campaigns [8]. Comic style narration often remains a widely accepted strategy due to their high potential to creatively engage the audience with the story and the message [20]. However, composing comic-style campaigns (CSC) requires dedicated and creative efforts from multiple experts throughout the process [1] such as story writer, script writer, graphic designer, visual artist, letterer, etc. Such setup is often costly and time-consuming. In this work, we present a novel human-AI collaborative framework (ComicCamp) that enables the CSC creators throughout the process to get crucial recommendations on creating story narration, aesthetic graphic illustration, and dialogues along with the relevant speech bubble layout across panels for effective narration of campaigns.




Figure 1: (a) The proposed collaborative framework ComicCamp for CSC with (A) Multi-Agent Narration Planner, (B) Graphics Illustrator, and (C) Letterer. (b) An example generated campaign based on the campaign description: *New way to manage your leaves with AI chatbot assistant. You can now easily apply for leaves, track leave request, and cancel leave request at one place..*

2 ComicCamp Framework

In this section, we present the collaborative framework ComicCamp for CSC creation. The framework has four main components (see Figure 1a):

1. **CSC creator:** The CSC creator is central to the ComicCamp framework. It could be a single individual or a team of creators who can leverage the recommendations from the multi-agent narration planner to visualize the narration with graphics illustrator and letterer.
2. **Multi-agent narration planner:** It is a foundational component in the ComicCamp framework as it is responsible for translating the campaign description to the narration plan using a multi-agent

collaboration between multiple LLM agents such as campaign manager, story composer, story critic, etc. These agents leverage the contextual information from the contextual knowledge base (CKB) which consist of the database of characters (and their personality traits), list of narrative structures [15, 18, 16, 21, 5], narration guidelines, types of speech bubbles and their description [12, 11, 10, 9], etc. For each panel, the generated narration plan consists of: setting, characters, speech bubble layouts, and dialogues (refer to Appendix for a detailed discussion).

3. **Graphics illustrator:** The CSC creator can adapt the setting and the characters recommendation from the narration plan to create the textual description for each panel. This description allows us to generate panels using text-to-image generation models such as stable diffusion [13]. The story is illustrated with a pre-defined set of characters and hence we need consistent characters across panels. We personalize the stable diffusion model by adapting the DreamBooth technique [14] to maintain multiple character consistency across panels while providing the visual reference of the characters to fine-tune the stable diffusion model (refer to Appendix for a detailed discussion).
4. **Letterer:** Finally, the CSC creator adapt the speech bubble layouts and the dialogues recommendation in the narrative plan to place them appropriately on each of the generated panel by the graphics illustrator.

3 A Case Study

In this section, we present a case study by applying the ComicCamp framework to internal broadcast campaigns for corporate communication in our organization. We consider existing campaigns covering diverse campaign topics such as webinar invitation, job posting, software upgrade, etc (see example campaigns in Figure 1b and Appendix). The CSC creator first brainstormed with GPT-4o [6, 3] to come up with the list of characters and their personality traits. We then synthetically generate these characters using FLUX.1 [4] text-to-image generation model. They have identified 10 unique comic characters in indie comic style and anime style each. We also consider 10 characters from comicgen [2] that are created by human artists. The visual reference of the characters is used by the graphic illustrator while the personality traits is used by the narration planner [19]. In total, the CSC creator creates 9 campaigns consisting of three styles (indie comic, anime, and comicgen). The CSC creator leverages the narration plan recommendation and makes necessary changes to story, setting, characters, and dialogues. In Table 1, we provide a summary of the degree of changes made by the CSC creator (as reported by them post campaign creation). We observe that the CSC creator made no changes to the story in 5 campaigns whereas they retained the recommendation for setting in 8 campaigns. They made relatively more changes to the recommended characters (2 low, 1 Medium, and 4 High) suggesting the complexity in picking the right character combination which is often preference driven. They also made limited changes to the suggested dialogues and never required a high degree of change across campaigns.

Campaign	Story	Setting	Characters	Dialogue
#1	Medium	No	Low	Low
#2	No	No	High	No
#3	No	No	High	Low
#4	No	No	No	Low
#5	No	No	No	No
#6	Low	Low	Medium	Medium
#7	Low	No	High	Low
#8	Medium	No	Low	Medium
#9	No	No	High	Low

Table 1: Degree of changes (No, Low, Medium, and High) made by the CSC creator to the narrative plan recommendation.

Campaign	Q1	Q2	Q3	Q4	Q5
#1	3.083	3.667	3.333	3.5	3.167
#2	3.917	3.667	3.75	3.917	3.833
#3	3.917	3.583	3.833	3.917	3.667
#4	3.833	3.583	3.583	3.583	3.583
#5	3.5	3.417	3.667	4	3.5
#6	4.083	3.917	3.833	4.167	4.25
#7	3.583	3.25	3.417	3.417	3.417
#8	3.833	3.75	3.833	3.833	3.833
#9	3.833	3.667	3.917	3.75	3.75
Overall	3.731	3.611	3.685	3.787	3.667

Table 2: Feedback from the users. We report the average score.

We conducted an internal survey with volunteers who have already received the original campaigns. We gave all the created campaigns to the evaluators along with the campaign description and the original campaign. For each campaign, we ask them five questions and asked them to provide a rating from 1 (very poor) to 5 (excellent): alignment with the description and the original campaign (Q1), coherence of the story (Q2), appropriateness of dialogues (Q3), connect between visual and textual narration (Q4), and overall quality (Q5). In total, 12 people participated in the survey (see Table 2). We observe that the average rating for Q1 to Q5 is 3.5 or higher for 7, 7, 7, 8, and 7 campaigns respectively. It further highlights the importance of our framework in collaboratively creating high-quality comic style campaigns with coherent story, high alignment with campaign description, and appropriate visuals and dialogues.

References

- [1] Overview of the Comic Creation Process. <https://comiconverse.com/overview-of-the-comic-creation-process-25482>, . Accessed: 2024-08-30.
- [2] Make your own comic. <https://gramener.com/comicgen/>, . Accessed: 2024-08-30.
- [3] Hello gpt-4o. <https://openai.com/index/hello-gpt-4o/>. Accessed: 2024-09-5.
- [4] Announcing black forest labs. <https://blackforestlabs.ai/announcing-black-forest-labs/>. Accessed: 2024-08-30.
- [5] Story Structure: 7 Types All Writers Should Know. <https://blog.reedsy.com/guide/story-structure/>. Accessed: 2024-08-30.
- [6] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- [7] Zirthing Lian Bawm and Rudra Pratap Deb Nath. A conceptual model for effective email marketing. In *2014 17th International Conference on Computer and Information Technology (ICCIT)*, pages 250–256. IEEE, 2014.
- [8] Julián Chaparro-Peláez, Ángel Hernández-García, and Ángel-José Lorente-Páramo. May i have your attention, please? an investigation on opening effectiveness in e-mail marketing. *Review of Managerial Science*, 16(7):2261–2284, 2022.
- [9] Neil Cohn. A visual lexicon. *Public Journal of Semiotics*, 1(1):35–56, 2007.
- [10] Neil Cohn. Comics, linguistics, and visual language: The past and future of a field. In *Linguistics and the Study of Comics*, pages 92–118. Springer, 2012.
- [11] Neil Cohn. Beyond speech balloons and thought bubbles: The integration of text and image. *Semiotica*, 2013(197):35–63, 2013.
- [12] Neil Cohn. *The Visual Language of Comics: Introduction to the Structure and Cognition of Sequential Images*. A&C Black, 2013.
- [13] Robin Rombach, Andreas Blattmann, Dominik Lorenz, Patrick Esser, and Björn Ommer. High-resolution image synthesis with latent diffusion models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 10684–10695, 2022.
- [14] Nataniel Ruiz, Yuanzhen Li, Varun Jampani, Yael Pritch, Michael Rubinstein, and Kfir Aberman. Dreambooth: Fine tuning text-to-image diffusion models for subject-driven generation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pages 22500–22510, 2023.
- [15] Tammy A Sincavage. Effectively structuring rising action in fiction: Using “fichte’s curve” to build a sturdy story toward a worthy climax.
- [16] Zihan Song, Robert E Roth, Lily Houtman, Timothy Prestby, Alicia Iverson, and Song Gao. Visual storytelling with maps: An empirical study on story map themes and narrative elements, visual storytelling genres and tropes, and individual audience differences. *cartographic perspectives*, (100):10–44, 2022.
- [17] Beidou Wang, Martin Ester, Yikang Liao, Jiajun Bu, Yu Zhu, Ziyu Guan, and Deng Cai. The million domain challenge: Broadcast email prioritization by cross-domain recommendation. In *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 1895–1904, 2016.
- [18] David Wilson. Using design knowledge to tell story. Master’s thesis, Kent State University, 2022.

- [19] Qingyun Wu, Gagan Bansal, Jieyu Zhang, Yiran Wu, Beibin Li, Erkang Zhu, Li Jiang, Xiaoyun Zhang, Shaokun Zhang, Jiale Liu, et al. Autogen: Enabling next-gen llm applications via multi-agent conversation. In *ICLR 2024 Workshop on Large Language Model (LLM) Agents*.
- [20] Ziang Xiao, Po-Shiun Ho, Xinran Wang, Karrie Karahalios, and Hari Sundaram. Should we use an abstract comic form to persuade? experiments with online charitable donation. *Proceedings of the ACM on Human-Computer Interaction*, 3(CSCW):1–28, 2019.
- [21] Leni Yang, Xian Xu, XingYu Lan, Ziyang Liu, Shunan Guo, Yang Shi, Huamin Qu, and Nan Cao. A design space for applying the freytag’s pyramid structure to data stories. *IEEE Transactions on Visualization and Computer Graphics*, 28(1):922–932, 2021.